

LFM2: Liquid Foundation Model 2

Step-by-Step Implementation Guide

Complete Tutorial

September 6, 2025

Contents

Contents	1
1 Introduction	2
2 Foundation Components	2
2.1 Configuration and Setup	2
2.2 RMSNorm Implementation	2
2.3 Rotary Positional Embedding (RoPE)	3
3 LFM2 Convolution Block	4
3.1 Architecture Diagram	4
3.2 Implementation	4
4 Grouped Query Attention	7
4.1 Architecture Diagram	7
4.2 Implementation	7
5 SwiGLU Feed-Forward Network	10
5.1 Implementation	10
6 LFM2 Transformer Block	11
6.1 Architecture Diagram	11
6.2 Implementation	11
7 Complete LFM2 Model	14
7.1 Architecture Diagram	14
7.2 Implementation	15
7.3 Language Modeling Head	18
7.4 Attention Mask Preparation	18
7.5 Model Creation and Testing	20
8 Conclusion	21

1 Introduction

LFM2 (Liquid Foundation Model 2) represents a hybrid neural architecture combining convolution operations for local processing with attention mechanisms for global dependencies. This tutorial walks through each component step-by-step, showing the architecture diagram first, followed immediately by the corresponding implementation.

2 Foundation Components

2.1 Configuration and Setup

Let's start with the basic configuration that defines our model:

```

1 import math
2 import torch
3 import torch.nn as nn
4 import torch.nn.functional as F
5 from typing import Optional, Tuple, Dict, Any, Union
6 from dataclasses import dataclass
7
8 @dataclass
9 class LFM2Config:
10     """Configuration class for LFM2 model."""
11     vocab_size: int = 32768          # Size of the vocabulary
12     hidden_size: int = 1024         # Hidden dimension size
13     intermediate_size: int = 2816   # Feedforward intermediate dimension
14     num_conv_blocks: int = 10       # Number of LIV convolution blocks
15     num_attention_blocks: int = 6   # Number of GQA blocks
16     num_attention_heads: int = 16   # Number of attention heads
17     num_key_value_heads: int = 4    # Number of key-value heads for GQA
18     conv_kernel_size: int = 3      # Kernel size for short convolutions
19     max_position_embeddings: int = 32768 # Maximum sequence length
20     rms_norm_eps: float = 1e-6     # RMS normalization epsilon
21     rope_theta: float = 10000.0    # RoPE base frequency
22     attention_dropout: float = 0.0 # Dropout rate for attention
23     hidden_dropout: float = 0.0    # Dropout rate for hidden layers
24     initializer_range: float = 0.02 # Weight initialization std
25     use_cache: bool = True          # Enable KV caching
26     tie_word_embeddings: bool = True # Tie input/output embeddings

```

This configuration defines the hybrid architecture: 10 convolution blocks for local processing, followed by 6 attention blocks for global dependencies. The grouped query attention uses 16 query heads but only 4 key-value heads (4:1 ratio) for memory efficiency.

2.2 RMSNorm Implementation

```

1 class RMSNorm(nn.Module):
2     """Root Mean Square Layer Normalization.
3
4     RMSNorm normalizes using only the root mean square, without centering:
5     RMSNorm(x) = (x / RMS(x)) * learnable_scale
6     where RMS(x) = sqrt(mean(x^2))
7     """
8

```

```

9  def __init__(self, hidden_size: int, eps: float = 1e-6):
10     super().__init__()
11     # Learnable scale parameter (no bias unlike LayerNorm)
12     self.weight = nn.Parameter(torch.ones(hidden_size))
13     self.variance_epsilon = eps
14
15     def forward(self, hidden_states: torch.Tensor) -> torch.Tensor:
16         # Store original dtype for final output
17         input_dtype = hidden_states.dtype
18
19         # Convert to float32 for stable computation
20         hidden_states = hidden_states.to(torch.float32)
21
22         # Compute variance: mean of squared values
23         variance = hidden_states.pow(2).mean(-1, keepdim=True)
24
25         # Normalize by RMS:  $x / \sqrt{\text{variance} + \text{eps}}$ 
26         hidden_states = hidden_states * torch.rsqrt(variance + self.
27             variance_epsilon)
28
29         # Apply learnable scale and convert back to original dtype
30         return self.weight * hidden_states.to(input_dtype)

```

RMSNorm is more stable than LayerNorm because it doesn't center the data (no mean subtraction), only scales by the RMS. This reduces computation and often improves training stability.

2.3 Rotary Positional Embedding (RoPE)

```

1  class RotaryPositionalEmbedding(nn.Module):
2      """Rotary Positional Embedding implementation.
3
4      RoPE encodes position by rotating query and key vectors by position-dependent
5      angles.
6      This allows the model to understand relative positions naturally.
7      """
8
9      def __init__(self, dim: int, max_position_embeddings: int = 32768,
10         base: float = 10000.0):
11         super().__init__()
12         self.dim = dim
13         self.max_position_embeddings = max_position_embeddings
14         self.base = base
15
16         # Precompute frequency inverse:  $1 / (\text{base}^{(2i/d)})$ 
17         # This creates different rotation frequencies for each dimension pair
18         inv_freq = 1.0 / (self.base **
19             (torch.arange(0, self.dim, 2).float() / self.dim))
20         self.register_buffer("inv_freq", inv_freq, persistent=False)
21
22     def forward(self, x: torch.Tensor, seq_len: int) -> Tuple[torch.Tensor, torch.
23         Tensor]:
24         # Generate position indices [0, 1, 2, ..., seq_len-1]
25         t = torch.arange(seq_len, device=x.device).type_as(self.inv_freq)
26
27         # Compute frequencies for each position: pos * inv_freq
28         freqs = torch.outer(t, self.inv_freq)

```

```

28     # Duplicate frequencies for sine and cosine (each dim pair needs both)
29     emb = torch.cat((freqs, freqs), dim=-1)
30
31     # Return cosine and sine components
32     cos = emb.cos()
33     sin = emb.sin()
34
35     return cos, sin
36
37 def rotate_half(x: torch.Tensor) -> torch.Tensor:
38     """Rotates half the hidden dims of the input.
39
40     For RoPE, we need to rotate pairs of dimensions. This function
41     swaps the first and second half and negates the second half.
42     """
43     x1 = x[..., : x.shape[-1] // 2] # First half
44     x2 = x[..., x.shape[-1] // 2 :] # Second half
45     return torch.cat((-x2, x1), dim=-1) # Rotate: [-x2, x1]
46
47 def apply_rotary_pos_emb(q: torch.Tensor, k: torch.Tensor,
48                           cos: torch.Tensor, sin: torch.Tensor) -> Tuple[torch.
49                               Tensor, torch.Tensor]:
50     """Apply rotary positional embedding to query and key tensors.
51
52     The rotation is: x_rotated = x * cos + rotate_half(x) * sin
53     This applies the rotation matrix in complex number form.
54     """
55     q_embed = (q * cos) + (rotate_half(q) * sin)
56     k_embed = (k * cos) + (rotate_half(k) * sin)
57     return q_embed, k_embed

```

RoPE works by treating pairs of dimensions as complex numbers and rotating them by position-dependent angles. This creates a natural way for the model to understand relative positions between tokens.

3 LFM2 Convolution Block

3.1 Architecture Diagram

3.2 Implementation

```

1 class LFM2ConvBlock(nn.Module):
2     """LFM2 double-gated short-range convolution block.
3
4     This implements Linear Input-Varying (LIV) convolution:
5     1. Project input to 3 components: gates B, C, and values x'
6     2. First gating: B * x' (input-dependent gating)
7     3. Depthwise convolution for local pattern extraction
8     4. Second gating: C * conv_output (input-dependent gating)
9     5. Final projection back to hidden size
10    """
11
12    def __init__(self, config: LFM2Config):
13        super().__init__()
14        self.config = config

```

LFM2 Convolution Block

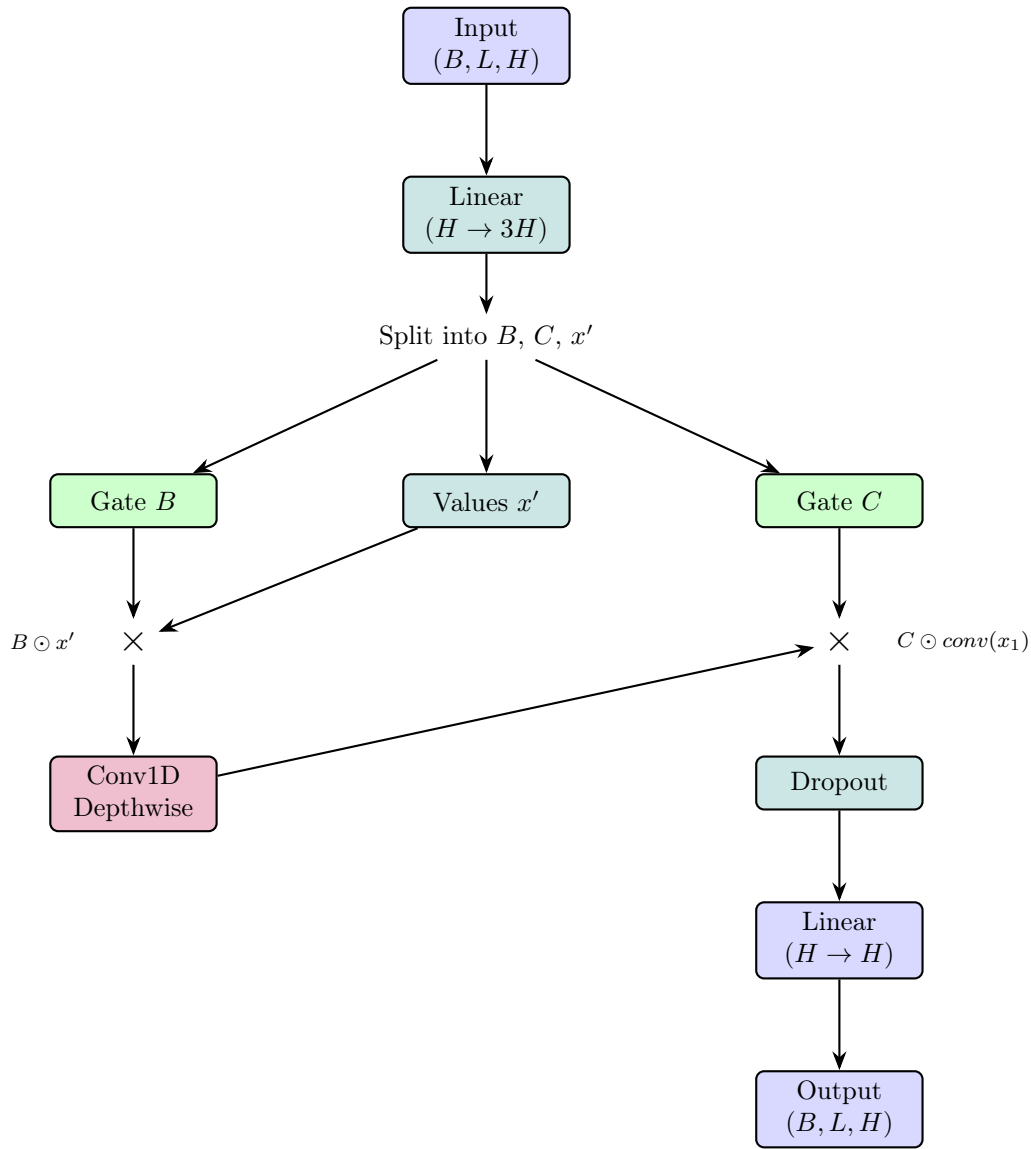


Figure 1: Double-gated Linear Input-Varying convolution with input-dependent gates B and C.

```

15     self.hidden_size = config.hidden_size
16
17     # Input projection to gates and values (splits into B, C, x)
18     self.input_projection = nn.Linear(
19         self.hidden_size,
20         3 * self.hidden_size, # Output: 3 * hidden_size for B, C, x
21         bias=False,
22     )
23
24     # Depthwise convolution: each channel processed independently
25     # This is memory efficient and captures local patterns

```

```

26     self.conv = nn.Conv1d(
27         self.hidden_size,          # Input channels
28         self.hidden_size,          # Output channels (same as input)
29         kernel_size=config.conv_kernel_size, # Usually 3
30         padding=config.conv_kernel_size // 2, # Same padding
31         groups=self.hidden_size,      # Depthwise: groups = channels
32         bias=False,
33     )
34
35     # Output projection back to original hidden size
36     self.output_projection = nn.Linear(
37         self.hidden_size, self.hidden_size, bias=False
38     )
39
40     # Dropout for regularization
41     self.dropout = nn.Dropout(config.hidden_dropout)
42
43 def forward(self, x: torch.Tensor) -> torch.Tensor:
44     """Forward pass implementing the LIV convolution.
45
46     Args:
47         x: Input tensor of shape (batch_size, seq_len, hidden_size)
48
49     Returns:
50         Output tensor of same shape as input
51     """
52     batch_size, seq_len, hidden_size = x.shape
53
54     # Step 1: Project input to gates B, C and values x'
55     # Input: (B, L, H) -> Output: (B, L, 3H)
56     projected = self.input_projection(x)
57
58     # Split into three equal parts: B, C, x'
59     # Each has shape (B, L, H)
60     B, C, x_proj = projected.chunk(3, dim=-1)
61
62     # Step 2: First gating - multiply values by gate B
63     # This allows input-dependent filtering before convolution
64     x_gated = B * x_proj # Element-wise multiplication
65
66     # Step 3: Apply depthwise convolution
67     # Conv1D expects (batch, channels, sequence), so transpose
68     x_conv_input = x_gated.transpose(1, 2) # (B, L, H) -> (B, H, L)
69     x_conv = self.conv(x_conv_input) # Apply convolution
70     x_conv = x_conv.transpose(1, 2) # Back to (B, L, H)
71
72     # Step 4: Second gating - multiply conv output by gate C
73     # This provides input-dependent filtering after convolution
74     x_gated_2 = C * x_conv
75
76     # Step 5: Apply dropout for regularization
77     x_gated_2 = self.dropout(x_gated_2)
78
79     # Step 6: Final output projection
80     output = self.output_projection(x_gated_2)
81
82     return output

```

The key innovation here is the double gating mechanism. Unlike standard convolution, both

gates B and C depend on the input, making this a "Linear Input-Varying" convolution. The depth-wise convolution ($\text{groups}=\text{hidden_size}$) is computationally efficient while still capturing local patterns.

4 Grouped Query Attention

4.1 Architecture Diagram

4.2 Implementation

```

1 class GroupedQueryAttention(nn.Module):
2     """Grouped Query Attention implementation.
3
4     GQA reduces memory usage by using fewer key-value heads than query heads.
5     For example: 16 query heads, 4 key-value heads means each KV head
6     is shared across 4 query heads (16/4 = 4 groups).
7     """
8
9     def __init__(self, config: LFM2Config):
10         super().__init__()
11         self.config = config
12         self.hidden_size = config.hidden_size
13         self.num_heads = config.num_attention_heads # e.g., 16
14         self.num_key_value_heads = config.num_key_value_heads # e.g., 4
15         self.head_dim = self.hidden_size // self.num_heads # e.g., 64
16
17         # Number of query heads per key-value head
18         self.num_key_value_groups = self.num_heads // self.num_key_value_heads
19
20         # Validate configuration
21         if self.hidden_size % self.num_heads != 0:
22             raise ValueError(f"hidden_size must be divisible by num_heads")
23
24         # Linear projections - note different output sizes
25         self.q_proj = nn.Linear(
26             self.hidden_size,
27             self.num_heads * self.head_dim, # Full size: 16 * 64
28             bias=False,
29         )
30         self.k_proj = nn.Linear(
31             self.hidden_size,
32             self.num_key_value_heads * self.head_dim, # Reduced: 4 * 64
33             bias=False,
34         )
35         self.v_proj = nn.Linear(
36             self.hidden_size,
37             self.num_key_value_heads * self.head_dim, # Reduced: 4 * 64
38             bias=False,
39         )
40         self.o_proj = nn.Linear(
41             self.num_heads * self.head_dim, # Back to full hidden size
42             self.hidden_size,
43             bias=False,
44         )
45
46         # Rotary positional embedding
47         self.rotary_emb = RotaryPositionalEmbedding(

```

Grouped Query Attention

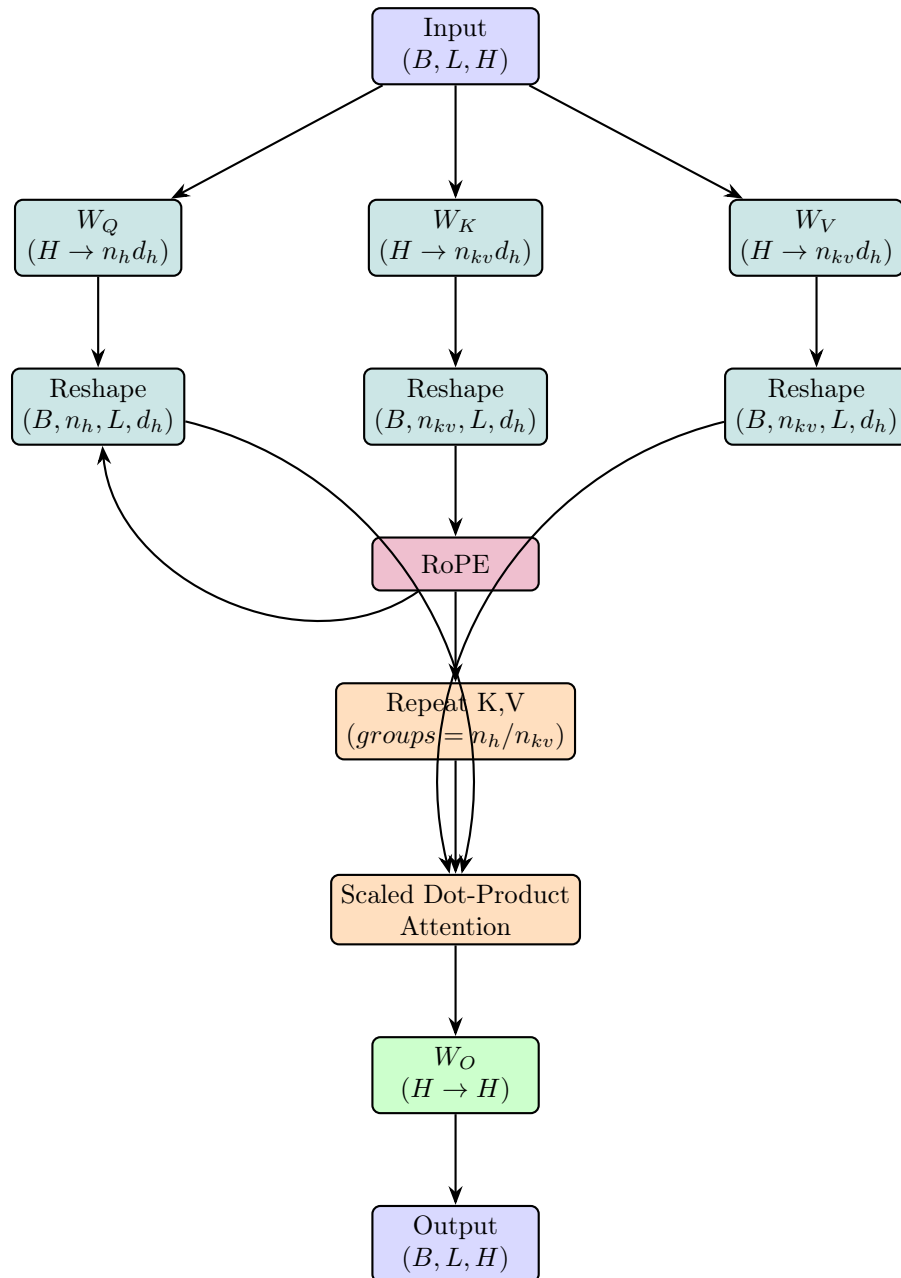


Figure 2: GQA uses fewer key-value heads than query heads, then repeats K,V to match Q heads.

```

48         self.head_dim,
49         config.max_position_embeddings,
50         config.rope_theta,
51     )
52
53     # Attention dropout
54     self.attention_dropout = nn.Dropout(config.attention_dropout)

```



```

55
56     def forward(self, hidden_states: torch.Tensor,
57                 attention_mask: Optional[torch.Tensor] = None,
58                 position_ids: Optional[torch.Tensor] = None,
59                 past_key_value: Optional[Tuple[torch.Tensor]] = None,
60                 output_attentions: bool = False,
61                 use_cache: bool = False) -> Tuple[torch.Tensor, Optional[torch.
62                     Tensor], Optional[Tuple[torch.Tensor]]]:
63
64     bsz, q_len, _ = hidden_states.size()
65
66     # Step 1: Project input to Q, K, V
67     query_states = self.q_proj(hidden_states) # (B, L, n_h * d_h)
68     key_states = self.k_proj(hidden_states) # (B, L, n_kv * d_h)
69     value_states = self.v_proj(hidden_states) # (B, L, n_kv * d_h)
70
71     # Step 2: Reshape for multi-head attention
72     # Q: (B, L, n_h * d_h) -> (B, n_h, L, d_h)
73     query_states = query_states.view(
74         bsz, q_len, self.num_heads, self.head_dim
75     ).transpose(1, 2)
76
77     # K, V: (B, L, n_kv * d_h) -> (B, n_kv, L, d_h)
78     key_states = key_states.view(
79         bsz, q_len, self.num_key_value_heads, self.head_dim
80     ).transpose(1, 2)
81     value_states = value_states.view(
82         bsz, q_len, self.num_key_value_heads, self.head_dim
83     ).transpose(1, 2)
84
85     # Step 3: Get sequence length for RoPE (including past context)
86     kv_seq_len = key_states.shape[-2]
87     if past_key_value is not None:
88         kv_seq_len += past_key_value[0].shape[-2]
89
90     # Step 4: Apply rotary positional embedding to Q and K
91     cos, sin = self.rotary_emb(value_states, kv_seq_len)
92     query_states, key_states = apply_rotary_pos_emb(
93         query_states, key_states, cos, sin
94     )
95
96     # Step 5: Handle past key-value cache for generation
97     if past_key_value is not None:
98         # Concatenate past and current key-values
99         key_states = torch.cat([past_key_value[0], key_states], dim=2)
100         value_states = torch.cat([past_key_value[1], value_states], dim=2)
101
102     # Prepare cache for next iteration
103     past_key_value = (key_states, value_states) if use_cache else None
104
105     # Step 6: Repeat K, V to match number of query heads (GQA core)
106     # From (B, n_kv, L, d_h) to (B, n_h, L, d_h)
107     key_states = key_states.repeat_interleave(
108         self.num_key_value_groups, dim=1
109     )
110     value_states = value_states.repeat_interleave(
111         self.num_key_value_groups, dim=1
112     )

```

```

113     # Step 7: Compute scaled dot-product attention
114     # Q @ K^T / sqrt(d_k)
115     attn_weights = torch.matmul(
116         query_states, key_states.transpose(2, 3)
117     ) / math.sqrt(self.head_dim)
118
119     # Apply attention mask if provided (for padding/causality)
120     if attention_mask is not None:
121         attn_weights = attn_weights + attention_mask
122
123     # Step 8: Apply softmax and dropout
124     attn_weights = F.softmax(attn_weights, dim=-1, dtype=torch.float32)
125     attn_weights = attn_weights.to(query_states.dtype)
126     attn_weights = self.attention_dropout(attn_weights)
127
128     # Step 9: Apply attention to values
129     attn_output = torch.matmul(attn_weights, value_states)
130
131     # Step 10: Reshape back to original format and apply output projection
132     attn_output = attn_output.transpose(1, 2).contiguous()
133     attn_output = attn_output.reshape(bsz, q_len, self.hidden_size)
134     attn_output = self.o_proj(attn_output)
135
136     # Return outputs (optionally include attention weights)
137     if not output_attentions:
138         attn_weights = None
139
140     return attn_output, attn_weights, past_key_value

```

The memory efficiency of GQA comes from step 6: instead of computing separate K,V for each head, we compute fewer K,V heads and repeat them. This reduces the KV cache size significantly during generation.

5 SwiGLU Feed-Forward Network

5.1 Implementation

```

1 class SwiGLU(nn.Module):
2     """SwiGLU activation function.
3
4     SwiGLU combines Swish activation with Gated Linear Units:
5     SwiGLU(x) = Swish(xW_gate) * (xW_up) * W_down
6
7     Where Swish(x) = x * sigmoid(x) = x * (x)
8     This has been shown to work better than ReLU-based FFNs.
9     """
10
11     def __init__(self, config: LFM2Config):
12         super().__init__()
13         self.hidden_size = config.hidden_size
14         self.intermediate_size = config.intermediate_size # Usually ~2.7x
15             hidden_size
16
17         # Two parallel transformations to intermediate size
18         self.gate_proj = nn.Linear(
19             self.hidden_size, self.intermediate_size, bias=False

```

```

19         )
20         self.up_proj = nn.Linear(
21             self.hidden_size, self.intermediate_size, bias=False
22         )
23
24         # Down projection back to hidden size
25         self.down_proj = nn.Linear(
26             self.intermediate_size, self.hidden_size, bias=False
27         )
28
29     def forward(self, x: torch.Tensor) -> torch.Tensor:
30         """Forward pass implementing SwiGLU.
31
32         Args:
33             x: Input tensor of shape (... , hidden_size)
34
35         Returns:
36             Output tensor of same shape as input
37         """
38         # Apply both linear projections in parallel
39         gate = self.gate_proj(x)      # "Gate" branch: controls information flow
40         up = self.up_proj(x)          # "Up" branch: carries information
41
42         # Apply SiLU (Swish) activation to gate branch
43         # SiLU(x) = x * sigmoid(x) - smooth, non-monotonic activation
44         swish_gate = F.silu(gate)
45
46         # Multiply gate and up branches (gating mechanism)
47         gated_output = swish_gate * up
48
49         # Project back down to original hidden size
50         return self.down_proj(gated_output)

```

SwiGLU has been empirically shown to outperform ReLU-based feed-forward networks. The gating mechanism allows the network to control information flow, while SiLU provides smooth gradients compared to ReLU.

6 LFM2 Transformer Block

6.1 Architecture Diagram

6.2 Implementation

```

1 class LFM2Block(nn.Module):
2     """LFM2 transformer block with GQA and SwiGLU.
3
4     Uses pre-normalization pattern:
5     1. x = x + Attention(RMSNorm(x))
6     2. x = x + SwiGLU(RMSNorm(x))
7
8     Pre-norm helps with training stability compared to post-norm.
9     """
10
11     def __init__(self, config: LFM2Config):
12         super().__init__()

```

LFM2 Transformer Block




Figure 3: Pre-normalization transformer block with residual connections.

```

13     self.config = config
14     self.hidden_size = config.hidden_size
15
16     # Main components
17     self.self_attn = GroupedQueryAttention(config)
18     self.mlp = SwiGLU(config)
19
20     # Layer normalization (pre-norm pattern)
21     self.input_layernorm = RMSNorm(
22         config.hidden_size, eps=config.rms_norm_eps

```

```

23         )
24         self.post_attention_layernorm = RMSNorm(
25             config.hidden_size, eps=config.rms_norm_eps
26         )
27
28     def forward(self, hidden_states: torch.Tensor,
29                 attention_mask: Optional[torch.Tensor] = None,
30                 position_ids: Optional[torch.Tensor] = None,
31                 past_key_value: Optional[Tuple[torch.Tensor]] = None,
32                 output_attentions: Optional[bool] = False,
33                 use_cache: Optional[bool] = False) -> Tuple[torch.Tensor, Optional
34                     [torch.Tensor], Optional[Tuple[torch.Tensor]]]:
35         """Forward pass implementing pre-norm transformer block."""
36
37         # First residual block: Self-attention
38         residual = hidden_states # Store for residual connection
39
40         # Pre-normalization before attention
41         hidden_states = self.input_layernorm(hidden_states)
42
43         # Apply self-attention
44         hidden_states, self_attn_weights, present_key_value = self.self_attn(
45             hidden_states=hidden_states,
46             attention_mask=attention_mask,
47             position_ids=position_ids,
48             past_key_value=past_key_value,
49             output_attentions=output_attentions,
50             use_cache=use_cache,
51         )
52
53         # Add residual connection
54         hidden_states = residual + hidden_states
55
56         # Second residual block: Feed-forward network
57         residual = hidden_states # Store for next residual connection
58
59         # Pre-normalization before MLP
60         hidden_states = self.post_attention_layernorm(hidden_states)
61
62         # Apply SwiGLU feed-forward network
63         hidden_states = self.mlp(hidden_states)
64
65         # Add residual connection
66         hidden_states = residual + hidden_states
67
68         # Prepare outputs
69         outputs = (hidden_states,)
70
71         if output_attentions:
72             outputs += (self_attn_weights,)
73
74         if use_cache:
75             outputs += (present_key_value,)
76
77         return outputs

```

The pre-normalization pattern (norm before operation) has become standard in modern transformers because it provides more stable gradients during training compared to post-normalization.

7 Complete LFM2 Model

7.1 Architecture Diagram

LFM2 (Liquid Foundation Model 2)

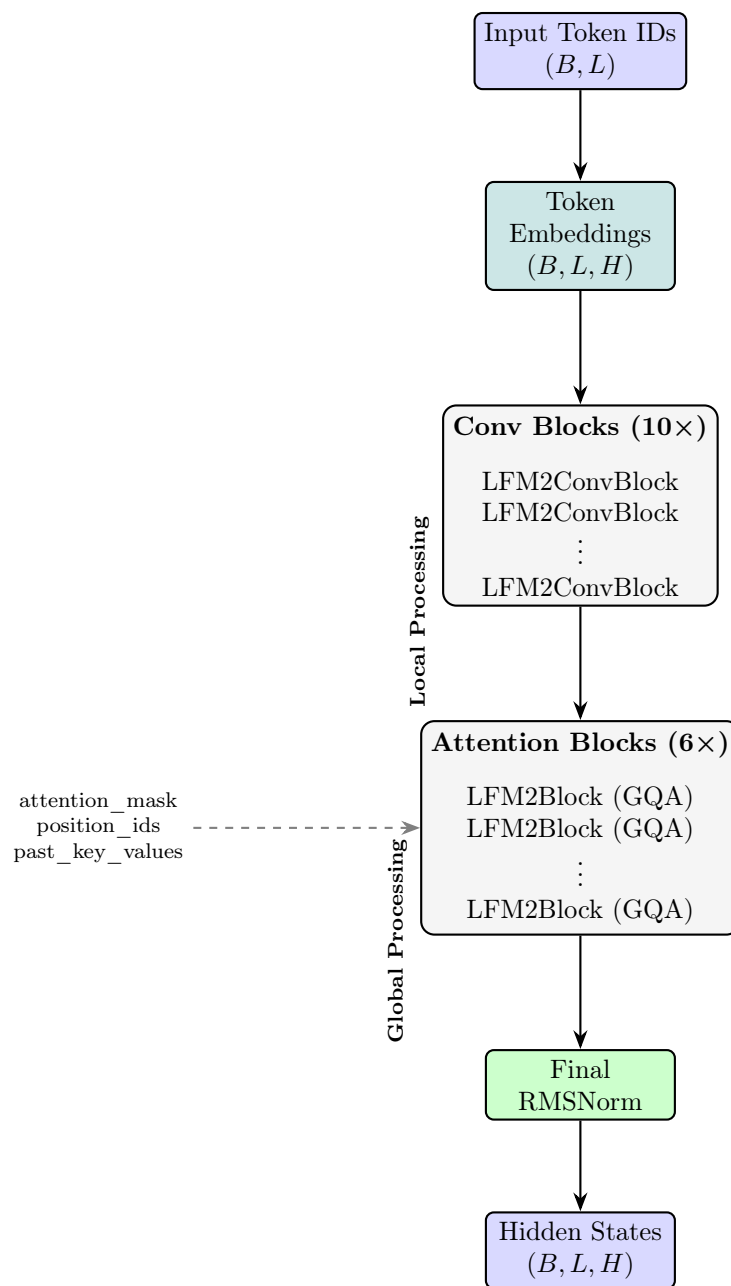


Figure 4: Complete LFM2 hybrid architecture: convolution for local patterns, attention for global dependencies.

7.2 Implementation

```

1 class LFM2Model(nn.Module):
2     """Complete LFM2 model with hybrid architecture.
3
4     Architecture:
5     - Token embeddings
6     - 10 LFM2ConvBlocks (local pattern processing)
7     - 6 LFM2Blocks with GQA (global dependency modeling)
8     - Final RMSNorm
9
10    This hybrid approach allows efficient local processing followed
11    by global attention for long-range dependencies.
12    """
13
14    def __init__(self, config: LFM2Config):
15        super().__init__()
16        self.config = config
17        self.vocab_size = config.vocab_size
18        self.hidden_size = config.hidden_size
19
20        # Token embeddings: convert token IDs to dense vectors
21        self.embed_tokens = nn.Embedding(config.vocab_size, config.hidden_size)
22
23        # Model layers: hybrid convolution + attention
24        self.layers = nn.ModuleList()
25
26        # Stage 1: Add convolution blocks (local processing)
27        for i in range(config.num_conv_blocks):
28            self.layers.append(LFM2ConvBlock(config))
29
30        # Stage 2: Add attention blocks (global processing)
31        for i in range(config.num_attention_blocks):
32            self.layers.append(LFM2Block(config))
33
34        # Final layer normalization
35        self.norm = RMSNorm(config.hidden_size, eps=config.rms_norm_eps)
36
37        # Initialize weights
38        self.apply(self._init_weights)
39
40    def _init_weights(self, module: nn.Module) -> None:
41        """Initialize model weights following standard practices."""
42        if isinstance(module, nn.Linear):
43            # Normal initialization for linear layers
44            torch.nn.init.normal_(
45                module.weight, mean=0.0, std=self.config.initializer_range
46            )
47            if module.bias is not None:
48                torch.nn.init.zeros_(module.bias)
49        elif isinstance(module, nn.Embedding):
50            # Normal initialization for embeddings
51            torch.nn.init.normal_(
52                module.weight, mean=0.0, std=self.config.initializer_range
53            )
54        elif isinstance(module, nn.Conv1d):
55            # Normal initialization for convolutions
56            torch.nn.init.normal_(

```

```

57         module.weight, mean=0.0, std=self.config.initializer_range
58     )
59     if module.bias is not None:
60         torch.nn.init.zeros_(module.bias)
61
62     def forward(self, input_ids: torch.LongTensor,
63                 attention_mask: Optional[torch.Tensor] = None,
64                 position_ids: Optional[torch.LongTensor] = None,
65                 past_key_values: Optional[Tuple[Tuple[torch.Tensor]]] = None,
66                 use_cache: Optional[bool] = None,
67                 output_hidden_states: Optional[bool] = None,
68                 return_dict: Optional[bool] = None) -> Union[Tuple, Dict[str,
69                     torch.Tensor]]:
70
71         """Forward pass through the complete LFM2 model."""
72
73         # Set default values
74         use_cache = use_cache if use_cache is not None else self.config.use_cache
75         output_hidden_states = (output_hidden_states if output_hidden_states is
76             not None
77             else self.config.output_hidden_states)
78         return_dict = return_dict if return_dict is not None else self.config.
79             use_return_dict
80
81         # Step 1: Convert token IDs to embeddings
82         hidden_states = self.embed_tokens(input_ids)
83         batch_size, seq_length = hidden_states.shape[:2]
84
85         # Step 2: Initialize position IDs if not provided
86         if position_ids is None:
87             device = input_ids.device if input_ids is not None else hidden_states.
88                 device
89             position_ids = torch.arange(seq_length, dtype=torch.long, device=
90                 device)
91             position_ids = position_ids.unsqueeze(0).expand(batch_size, -1)
92
93         # Step 3: Initialize past key values for caching
94         if past_key_values is None:
95             past_key_values = [None] * len(self.layers)
96
97         # Step 4: Prepare attention mask for causal modeling
98         if attention_mask is not None:
99             attention_mask = self._prepare_attention_mask(
100                 attention_mask, seq_length, batch_size
101             )
102
103         # Step 5: Initialize output containers
104         all_hidden_states = () if output_hidden_states else None
105         next_decoder_cache = () if use_cache else None
106
107         # Step 6: Process through all layers
108         for idx, decoder_layer in enumerate(self.layers):
109             if output_hidden_states:
110                 all_hidden_states += (hidden_states,)
111
112             past_key_value = past_key_values[idx] if past_key_values is not None
113                 else None
114
115             if isinstance(decoder_layer, LFM2ConvBlock):
116                 # Convolution blocks: simple forward pass (no attention cache)

```



```

110         hidden_states = decoder_layer(hidden_states)
111         layer_outputs = (hidden_states,)
112     else:
113         # Attention blocks: full forward pass with caching
114         layer_outputs = decoder_layer(
115             hidden_states,
116             attention_mask=attention_mask,
117             position_ids=position_ids,
118             past_key_value=past_key_value,
119             use_cache=use_cache,
120         )
121         hidden_states = layer_outputs[0]
122
123         # Collect cache from attention layers only
124         if use_cache and isinstance(decoder_layer, LFM2Block):
125             next_decoder_cache += (layer_outputs[2],)
126
127     # Step 7: Apply final normalization
128     hidden_states = self.norm(hidden_states)
129
130     # Add final hidden state to output
131     if output_hidden_states:
132         all_hidden_states += (hidden_states,)
133
134     # Prepare final cache
135     next_cache = next_decoder_cache if use_cache else None
136
137     # Return in requested format
138     if not return_dict:
139         return tuple(v for v in [hidden_states, next_cache, all_hidden_states]
140                             if v is not None)
141
142     return {
143         "last_hidden_state": hidden_states,
144         "past_key_values": next_cache,
145         "hidden_states": all_hidden_states,
146     }
147
148 def _prepare_attention_mask(self, attention_mask: torch.Tensor,
149                             seq_length: int, batch_size: int) -> torch.Tensor:
150     """Prepare causal attention mask combining padding and causality."""
151
152     # Create causal mask (lower triangular)
153     causal_mask = torch.triu(
154         torch.ones((seq_length, seq_length), dtype=torch.bool),
155         diagonal=1, # Upper triangle (including diagonal) = True (masked)
156     ).to(attention_mask.device)
157
158     # Expand attention mask to 4D: (batch, 1, seq, seq)
159     expanded_mask = (attention_mask[:, None, None, :])
160         .expand(batch_size, 1, seq_length, seq_length)
161         .to(torch.float32)
162
163     # Convert to mask format: 1.0 = masked, 0.0 = not masked
164     expanded_mask = 1.0 - expanded_mask
165     causal_float_mask = causal_mask.to(torch.float32)
166
167     # Combine padding and causal masks
168     combined_mask = expanded_mask + causal_float_mask[None, None, :, :]

```

```

168
169     # Convert to attention weights: 0 = attend, -inf = mask
170     return combined_mask.masked_fill(combined_mask > 0, float("-inf"))

```

The LFM2Model orchestrates the entire hybrid architecture. The key insight is processing through convolution blocks first for local patterns, then attention blocks for global dependencies.

7.3 Language Modeling Head

LFM2ForCausalLM - Language Modeling Head

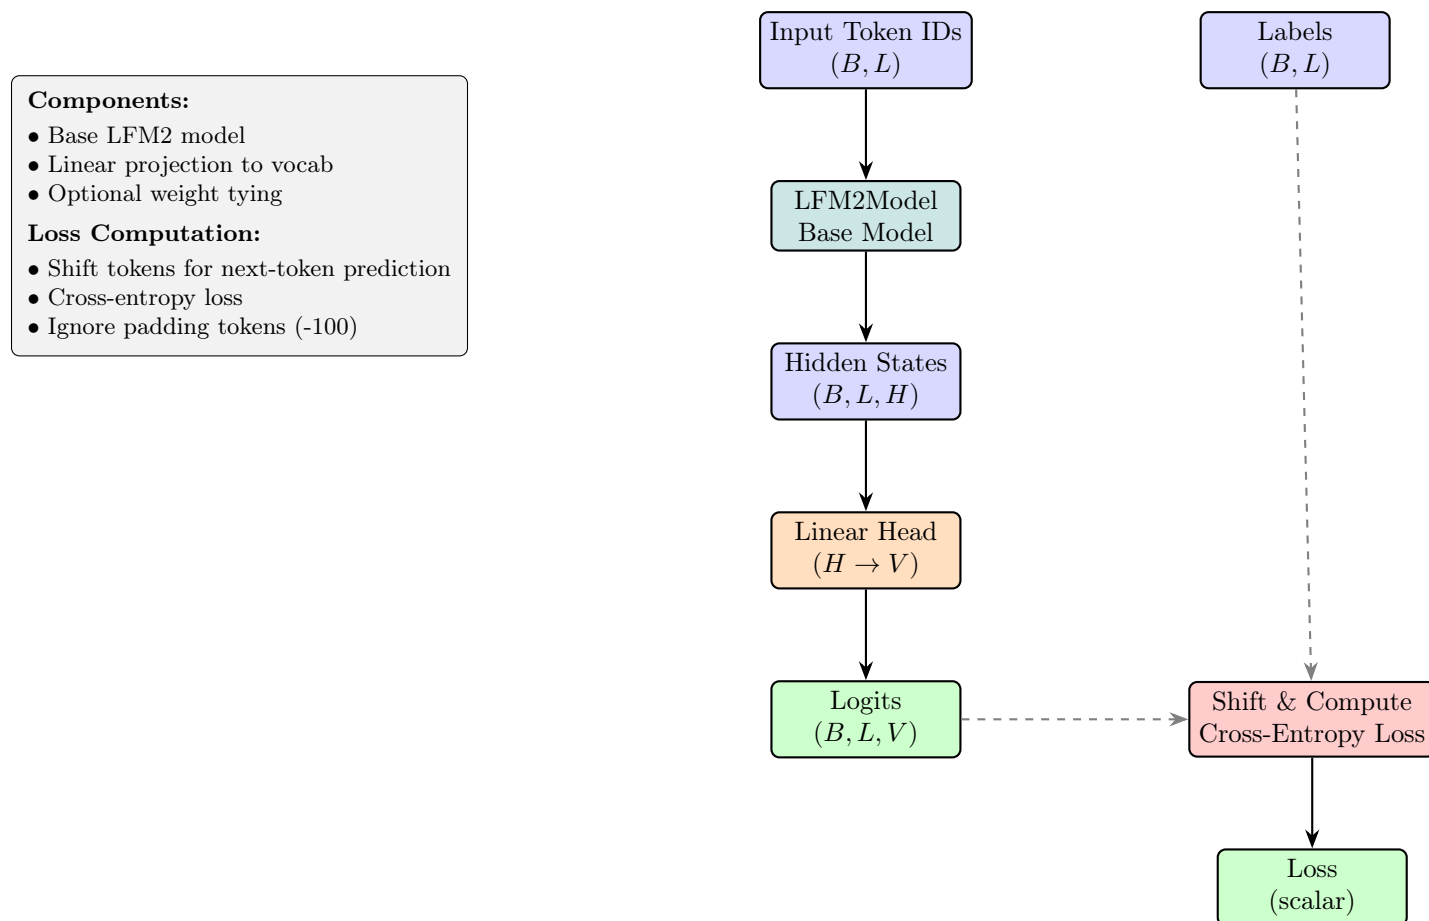


Figure 5: Language modeling head converts hidden states to vocabulary logits for text generation.

7.4 Attention Mask Preparation

```

1 class LFM2ForCausalLM(nn.Module):
2     """LFM2 model for causal language modeling (text generation)."""
3
4     def __init__(self, config: LFM2Config):

```

Attention Mask Preparation

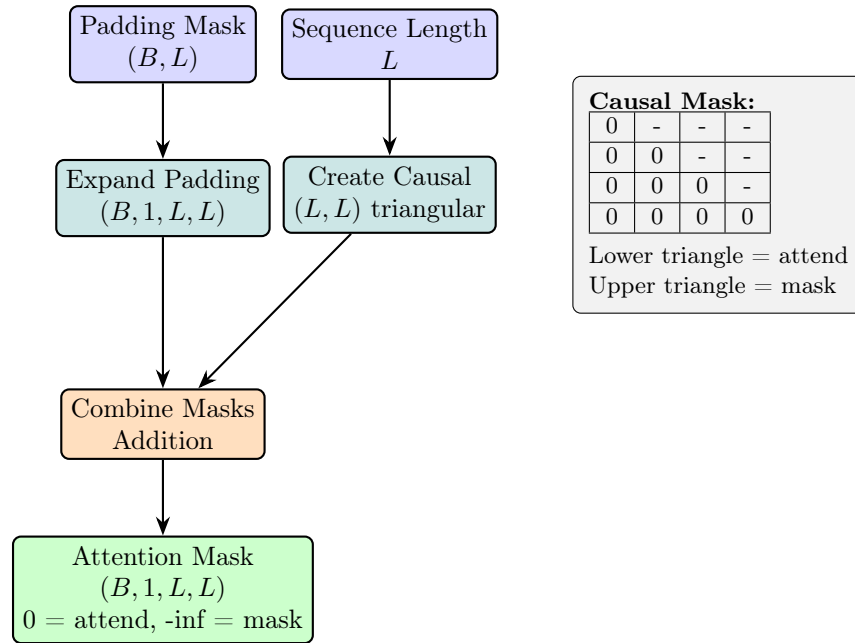


Figure 6: Attention mask combines padding mask with causal mask for autoregressive generation.

```

5     super().__init__()
6     self.config = config
7     self.model = LFM2Model(config)
8     self.vocab_size = config.vocab_size
9
10    # Language modeling head: hidden states -> vocabulary logits
11    self.lm_head = nn.Linear(config.hidden_size, config.vocab_size, bias=False
12                               )
13
14    # Optionally tie input and output embeddings (parameter sharing)
15    if config.tie_word_embeddings:
16        self.lm_head.weight = self.model.embed_tokens.weight
17
18    self.apply(self._init_weights)
19
20    def forward(self, input_ids: torch.LongTensor,
21                attention_mask: Optional[torch.Tensor] = None,
22                labels: Optional[torch.LongTensor] = None,
23                **kwargs) -> Union[Tuple, Dict[str, torch.Tensor]]:
24        """Forward pass for causal language modeling."""
25
26        # Forward through base model
27        outputs = self.model(
28            input_ids=input_ids,
29            attention_mask=attention_mask,
30            **kwargs
31        )

```

```

32     hidden_states = outputs["last_hidden_state"] if isinstance(outputs, dict)
33         else outputs[0]
34
35     # Compute logits for each position and vocabulary item
36     logits = self.lm_head(hidden_states)
37
38     # Compute loss if labels are provided (training)
39     loss = None
40     if labels is not None:
41         # Shift for next-token prediction: predict token i+1 from tokens 0..i
42         shift_logits = logits[..., :-1, :].contiguous() # Remove last
43             position
44         shift_labels = labels[..., 1:].contiguous() # Remove first
45             position
46
47         # Cross-entropy loss
48         loss_fct = nn.CrossEntropyLoss(ignore_index=-100)
49         shift_logits = shift_logits.view(-1, self.vocab_size)
50         shift_labels = shift_labels.view(-1)
51
52         # Move to same device for model parallelism
53         shift_labels = shift_labels.to(shift_logits.device)
54         loss = loss_fct(shift_logits, shift_labels)
55
56     # Return outputs
57     if isinstance(outputs, dict):
58         return {
59             "loss": loss,
60             "logits": logits,
61             "past_key_values": outputs.get("past_key_values"),
62             "hidden_states": outputs.get("hidden_states"),
63         }
64     else:
65         output = (logits,) + outputs[1:]
66         return (loss,) + output if loss is not None else output

```

7.5 Model Creation and Testing

```

1 def create_lfm2_model(model_size: str = "700M") -> LFM2ForCausalLM:
2     """Create LFM2 model with predefined configurations."""
3
4     # Predefined size configurations
5     size_configs = {
6         "350M": {
7             "hidden_size": 768,
8             "intermediate_size": 2048,
9             "num_attention_heads": 12,
10            "num_key_value_heads": 3,
11        },
12        "700M": {
13            "hidden_size": 1024,
14            "intermediate_size": 2816,
15            "num_attention_heads": 16,
16            "num_key_value_heads": 4,
17        },
18        "1.2B": {

```

```

19         "hidden_size": 1536,
20         "intermediate_size": 4096,
21         "num_attention_heads": 24,
22         "num_key_value_heads": 6,
23     },
24 }
25
26 # Create and configure model
27 config = LFM2Config()
28 config.__dict__.update(size_configs[model_size])
29
30 return LFM2ForCausalLM(config)
31
32 def test_lfm2_model():
33     """Test LFM2 model with dummy inputs."""
34     print("Creating LFM2 model...")
35     model = create_lfm2_model("350M")
36     model.eval()
37
38     # Create dummy inputs
39     batch_size, seq_len = 2, 128
40     input_ids = torch.randint(0, model.config.vocab_size, (batch_size, seq_len))
41     attention_mask = torch.ones(batch_size, seq_len)
42
43     print(f"Input shape: {input_ids.shape}")
44     print(f"Model parameters: {sum(p.numel() for p in model.parameters()):,}")
45
46     # Forward pass
47     with torch.no_grad():
48         outputs = model(input_ids=input_ids, attention_mask=attention_mask)
49
50     print(f"Output logits shape: {outputs['logits'].shape}")
51     print("Test completed successfully!")
52
53 # Run the test
54 test_lfm2_model()

```

8 Conclusion

This step-by-step implementation of LFM2 demonstrates several key innovations:

- **Hybrid Architecture:** Convolution blocks for local patterns + attention blocks for global dependencies
- **Grouped Query Attention:** Reduces memory usage by sharing key-value pairs across query heads
- **Modern Components:** RMSNorm, RoPE, SwiGLU for improved training and performance
- **Efficient Design:** Double-gated convolutions and depthwise operations for computational efficiency

The implementation provides a complete, working LFM2 model suitable for language modeling tasks, with careful attention to memory efficiency and modern best practices.